

Problem

It is hard to manage both, files and SQL data

Why?

It is hard to keep files and records in sync;

- Files get renamed or deleted
- New files appear
- No transactional semantics
- Different permission mechanisms
 - Unix vs. SQL users and permissions
- Different content access mechanisms
 - SQL connection vs. file access vs. web

Solution

Manipulating files with SQL

- Accessing contents of files
- ACID semantics
- Dealing with URLs
- File management (move, etc)

SQL/MED datalinks

according to SQL standard

- DATALINK datatype
Datatype for storing URLs
- Support functions
Using values of datalink type
- Special semantics
 - Transactions
 - Referential integrity
 - Access mediation
 - Point-in-time recovery

Other interesting libraries

- Built-in file functions – `pg_file_read...`
- Contrib adminpack – `pg_file_write`,
`pg_file_rename`, `pg_file_unlink...`
- PgFoundry uri – this could be replaced by
`url_part` and `url_part_set`
- CURL – it is best at grokking URLs. Talks
to complicated and secure web servers.
It could be used for actual web access.

Object hierarchies

Files

Volumes

«table»
dl_link

«table»
dl_prefix

Datalinks
Web, URLs

«table»
file_link

«table»
file_space

Files
Local disks
File paths

«table»
file_inode

«table»
file_dev

Devices, File content

APIs

«API»
datalink

SQL/MED functions

«API»
admin

Server administration functions

«API»
URL

URL utility functions

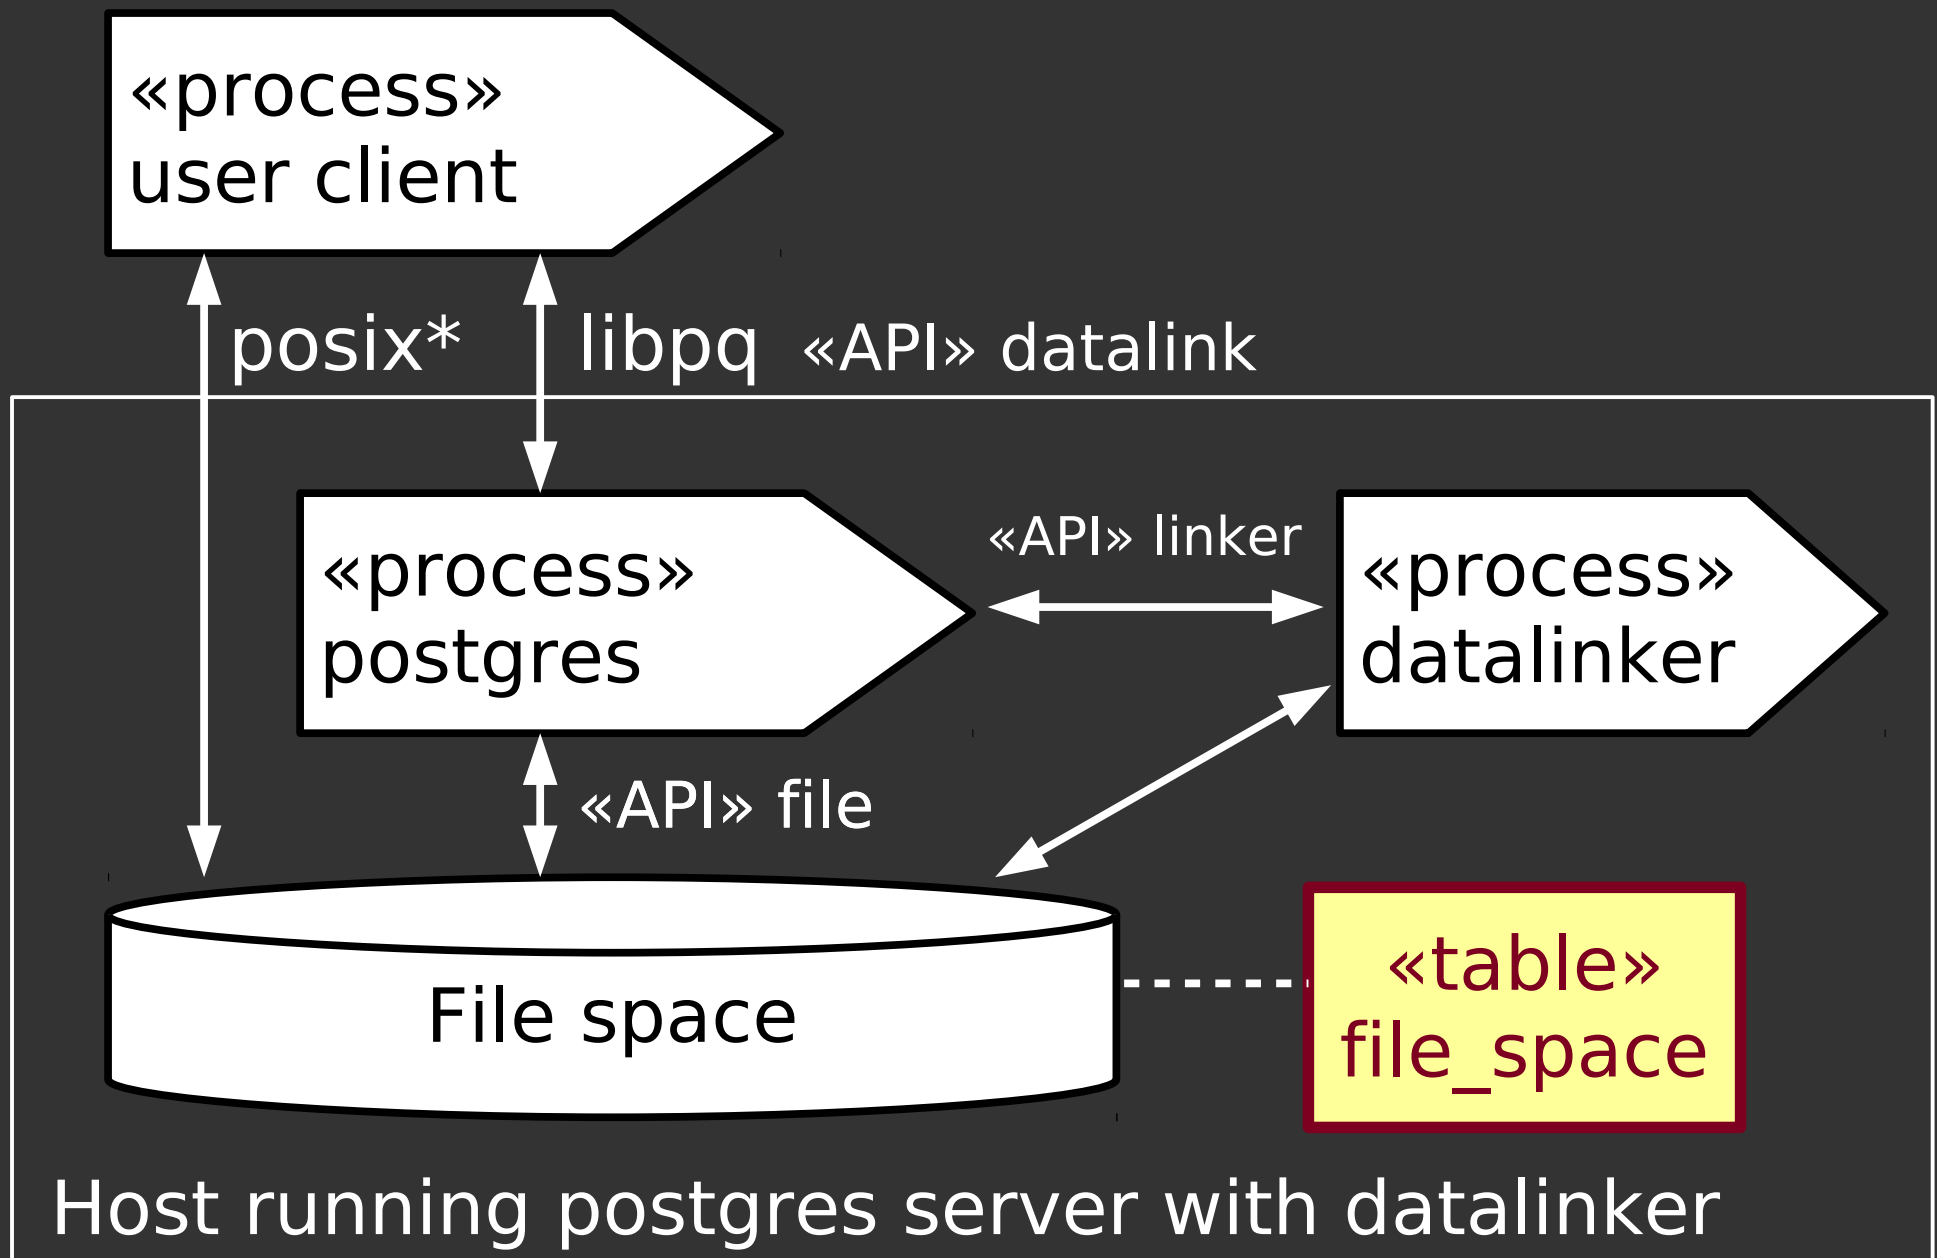
«API»
file

Transactional file functions

«API»
linker

Datalinker functions

Datalinker process



* only clients on same host

Datalinker responsibilities

- Backup a file
- Delete a file
- Link to SQL (protect) for reading
- Link to SQL (protect) for writing
- Checkpoint
- Rename a file (?)

Datalinker is accessed with «API» linker.
Not intended for end users!

«API» linker

- Internal datalink usage only
- Not intended for end users
- Stores instructions for datalinker in
«table» file_xact

«API» file – whole files

- Use `file_put()` to write whole files

```
SELECT file_put('/var/www/index.txt','Hello, world!');
```

- Use `file_get()` to read whole files

```
SELECT file_get('/var/www/index.txt');
```

```
Hello, world!
```

When writing files, new file is always created!
This is used to implement transactions.
Old files can be purged separately.

«API» file - incremental

- Use `file_begin()`, `file_write()` and `file_end()` to write files incrementally

This is useful for functions in procedural languages

```
BEGIN
  fh := file_begin('/var/www/index.html');
  FOR row IN
    SELECT label FROM items
  LOOP
    file_write(fh,row.label||E'\n')
  END LOOP;
  file_end(fh);
END
```